

GEORGIA INSTITUTE OF TECHNOLOGY  
SCHOOL of ELECTRICAL and COMPUTER ENGINEERING  
**ECE 2026      Fall 2013**  
**Lab #4: Synthesis of Sinusoidal Signals—Music Synthesis**

Date: 23 – 26 September 2013

---

Each Lab assignment in ECE2026 consists of three parts: Pre-Lab, In-lab Tasks, and Take-home Questions. It requires you to come into lab prepared. Be sure to read the entire lab carefully before arriving.

**Pre-Lab:** You should read the Pre-Lab section of the lab and go over all exercises in this section before going to your assigned lab session. Although you do not need to turn in results from the Pre-Lab, doing the exercises therein will help make your in-lab experience more rewarding and go more smoothly with less frustration and panicking.

**In-lab Tasks and Verification:** There are a number of designated tasks for each student to accomplish during the lab session. Students are encouraged to read, prepare for and even attempt at these tasks beforehand. These tasks must be completed **during your assigned lab time** and the steps marked *Instructor Verification* must also be signed off **during the lab time**. One of the laboratory instructors must verify the appropriate steps by signing on the **Instructor Verification** line. When you have completed a step that requires verification, simply put a plastic cup on top of your PC and demonstrate the step to one of the TAs or the professor. (You can also use the plastic cups to indicate if you have a more general question, i.e. you can use it to get our attention even if you don't have an Instructor Verification ready.)

**Take-home Questions:** At the end of each lab sheet below all the verification steps, several questions are to be answered by the student, who can choose to complete the answers while in the lab or after leaving the lab.

The lab sheet with all verification signatures and answers to the questions needs to be turned in to the Lab-grading TA at the beginning of the next lab session.

---

## 1 Introduction

This lab explores music/sound synthesis with sinusoidal waveforms of the form

$$x(t) = \sum_k A_k \cos(\omega_k t + \phi_k). \quad (1)$$

Since music consists of musical notes, it will be necessary to establish the connection between musical notes, their frequencies, and sinusoids. A secondary objective of the lab is to document the relationship between the synthesized signal, its spectrogram and the musical notes. The key take-home message is that simply adding sinusoids together can produce many interesting signals.

## 2 Pre-Lab

In this lab, the periodic waveforms and music signals will be created with the intention of playing them out through a speaker. Therefore, it is necessary to take into account the fact that a conversion is needed from the digital samples, which are numbers stored in the computer memory, to the actual voltage waveform that will be amplified and heard through the loudspeakers.

### 2.1 Theory of Sampling

Chapter 4 treats the subject of sampling in detail; we provide a quick summary of essential facts here. An idealized process of sampling a signal and the subsequent reconstruction of the signal from its samples is depicted in Fig. 1. This figure shows a continuous-time input signal  $x(t)$ , which is sampled by the continuous-to-discrete (C-to-D) converter to produce a sequence of samples  $x[n] = x(nT_s)$ , where  $n$  is the integer sample index and  $T_s$  is the sampling period. The sampling rate is  $f_s = 1/T_s$  where the units are samples per second. As described in Chapter 4 of the text, the ideal discrete-to-continuous (D-to-C) converter takes the input samples

and interpolates a smooth curve between them. The *Sampling Theorem* tells us that if the input signal  $x(t)$  is a sum of sinusoids, then the output  $y(t)$  will be equal to the input  $x(t)$  if the sampling rate is *more than twice the highest frequency*  $f_{\max}$  in the input, i.e.,  $f_s > 2f_{\max}$ . In other words, if we *sample fast enough* then there will be no problems synthesizing the continuous audio signal from  $x[n]$ . More rigorous expositions of the subject are covered in Chapter 4.




Figure 1: Sampling and reconstruction of a continuous-time signal.

## 2.2 A-to-D and D-to-A Conversion

Most computers have a built-in analog-to-digital (A-to-D) converter and a digital-to-analog (D-to-A) converter (usually on the sound card). These hardware systems are physical realizations of the idealized concepts of C-to-D and D-to-C converters, respectively. For purposes of this lab, nevertheless, we will assume that the hardware A/D and D/A represent indeed perfect realizations of C-to-D and D-to-C conversion.

The digital-to-analog conversion process has a number of aspects, but in its simplest form the only thing we need to worry about in this lab is that the time spacing  $T_s$  between the signal samples must correspond to the rate of the D-to-A hardware that is being used. From MATLAB, the sound output (converted from a numerical sequence to an acoustic signal) is done by the `soundsc(xx,fs)` function (and several other similar commands) which does support a variable D-to-A sampling rate if the hardware on the machine has such capability. A convenient choice for the D-to-A conversion rate is 11025 samples per second,<sup>1</sup> so  $T_s = 1/11025$  seconds; another common choice is 8000 samples/sec. Both of these rates satisfy the requirement of *sampling fast enough* as explained in the next section. In fact, most piano notes have relatively low frequencies, so an even lower sampling rate could be used. If you are using `soundsc()`, the vector `xx` will be scaled *automatically* for the D-to-A converter, but if you are using `sound.m`, you must scale the vector `xx` so that it lies between  $\pm 1$ . Consult `help sound`.

- (a) The ideal C-to-D converter is, in effect, being implemented whenever we take samples of a continuous-time formula, e.g.,  $x(t)$  at  $t = t_n$ . We do this in MATLAB by first making a vector of time indices, and then evaluating the formula for the continuous-time signal at the sample times, i.e.,  $x[n] = x(nT_s)$  if  $t_n = nT_s$ . This assumes perfect knowledge of the input signal, but we have already been doing it this way in previous labs.

To begin, create a vector `x1` of samples of a sinusoidal signal with  $A_1 = 100$ ,  $\omega_1 = 2\pi(800)$ , and  $\phi_1 = -\pi/3$ . Use a sampling rate of 11,025 samples/s, and compute a total number of samples equivalent to duration of 0.5 seconds. You may find it helpful to recall that a MATLAB statement such as `tt=0:0.01:3` would create a vector of numbers from 0 through 3 with increments of 0.01. Therefore, it is only necessary to determine the time increment needed to obtain 11025 samples in one second. You should use the `add_sines()` function from a previous lab for this part (with modification of the sampling rate).

Use `soundsc()` to play the resulting vector through the D-to-A converter of the computer, assuming that the hardware can support the  $f_s = 11025$  Hz rate. Listen to the output.

- (b) Now create another vector `x2` of samples of a second sinusoidal signal (0.8 secs in duration) for the case  $A_1 = 80$ ,  $\omega_1 = 2\pi(1200)$ , and  $\phi_1 = +\pi/4$ . Listen to the signal reconstructed from these samples. How does its sound compare to the signal in part (a)?
- (c) **Concatenate** the two signals `x1` and `x2` from the previous parts and put a short duration of 0.1 seconds of silence in-between. You should be able to concatenate by using a statement something like:

<sup>1</sup>This sampling rate is one quarter of the rate (44,100 Hz) used in audio CD players.

```
xx = [ x1, zeros(1,N), x2 ];
```

assuming that both **x1** and **x2** are row vectors. Determine the correct value of **N** to make 0.1 seconds of silence. Listen to this new signal to verify that it is correct.

- (d) To verify that the concatenation operation was done correctly in the previous part, make the following plot:

```
tt = (1/11025)*(1:length(xx)); plot( tt, xx );
```

This will plot a huge number of points, but it will show the “envelope” of the signal and verify that the amplitude changes from 100 to zero and then to 80 at the correct times. Notice that the time vector **tt** was created to have exactly the same length as the signal vector **xx**.

- (e) Now send the vector **xx** to the D-to-A converter again, but change the sampling rate **fs** parameter in **soundsc**(**xx**, **fs**) to 22050 samples/s. Do not re-compute the samples in **xx**, just tell the D-to-A converter that the sampling rate is 22050 samples/s. Describe how the *duration* and *pitch* of the signal were affected. Explain.

## 2.3 Structures in MATLAB

MATLAB can do structures. Structures are convenient for grouping information together. For example, run the following program which plots a sinusoid:

```
x.Amp = 7;
x.phase = -pi/2;
x.freq = 100;
x.fs = 11025
x.timeInterval = 0:(1/x.fs):0.05;
x.values = x.Amp*cos(2*pi*(x.freq)*(x.timeInterval) + x.phase);
x.name = 'My Signal';
x %---- echo the contents of the structure "x"
plot( x.timeInterval, x.values )
title( x.name )
```

Notice that the members of the structure can contain different types of variables: scalars, vectors or strings.

## 2.4 Debugging Skills

Testing and debugging code is a big part of any programming job, as you know if you’ve been staying up late on the first few labs. Almost any modern programming environment provides a *symbolic debugger* so that *break-points* can be set and variables examined in the middle of program execution. Nonetheless, many programmers still insist on using the old-fashioned method of inserting print statements in the middle of their code (or the MATLAB equivalent, leaving off a few semi-colons).

*In ECE2026 labs, it is your responsibility to debug your own codes.*

There are two ways to use debugging tools in MATLAB: via buttons in the edit window or via the command line. For help on the edit-window debugging features, access the menu **Help->Using the M-File Editor** which will pop up a browser window at the help page for editing and debugging. For a summary of the command-line debugging tools, try **help debug**. Here is part of what you’ll see:

```
dbstop - Set breakpoint.
dbclear - Remove breakpoint.
dbcont - Resume execution.
dbstack - List who called whom.
dbstatus - List all breakpoints.
dbstep - Execute one or more lines.
dbtype - List M-file with line numbers.
dbquit - Quit debug mode.
```

When a breakpoint is hit, MATLAB goes into debug mode. On the PC

and Macintosh the debugger window becomes active and on UNIX and VMS the prompt changes to a K>>. Any MATLAB command is allowed at the prompt. To resume M-file function execution, use DBCONT or DBSTEP. To exit from the debugger use DBQUIT.

One of the most useful modes of the debugger causes the program to jump into “debug mode” whenever an error occurs. This mode can be invoked by typing:

```
dbstop if error
```

With this mode active, you can snoop around inside a function and examine local variables that probably caused the error. You can also choose this option from the debugging menu in the MATLAB editor. It’s sort of like an automatic call to 911 when you’ve gotten into an accident. Try **help dbstop** for more information.

Download the file **coscos.m** and use the debugger to find the error(s) in the function. Call the function with the test case: **[xn,tn] = coscos(2,3,20,1)**. Use the debugger to:

1. Set a breakpoint to stop execution when an error occurs and jump into “Keyboard” mode;
2. Display the contents of important vectors while stopped;
3. Determine the size of all vectors by using either the **size()** function or the **whos** command; and
4. Lastly, modify variables while in the “Keyboard” mode of the debugger.

```
function [xx,tt] = coscos( f1, f2, fs, dur )
% COSCOS multiply two sinusoids
%
t1 = 0:(1/fs):dur;
t2 = 0:(1/f2):dur;
cos1 = cos(2*pi*f1*t1);
cos2 = cos(2*pi*f2*t2);
xx = cos1 .* cos2;
tt = t1;
```

## 2.5 Piano Keyboard

We’ll learn some music synthesis here.<sup>2</sup> (Section 4 involves synthesizing a music piece but is optional. If you choose to complete it, you will find it interesting and the experience rewarding.) Music signals use sinusoidal tones to produce audible sounds, generally known as music scales. A ready and intuitive mapping of a music scale to sound frequencies is the piano keyboard. A quick introduction to the layout of the piano keyboard is thus helpful. On a piano, the keyboard is divided into octaves—the notes in one octave being twice the frequency of the notes in the next lower octave. The white keys in each octave are named A through G. In order to define the frequencies of all the keys, one key must be designated as the reference. Usually, the reference note is the A above middle-C, called A-440 (or  $A_4$ ) because its frequency is 440 Hz.<sup>3</sup> Each octave contains 12 notes (5 black keys and 7 white) and the ratio between the frequencies of the notes is constant between successive notes. As a result, this ratio must be  $2^{1/12}$ . Since middle C is 9 keys below A-440, its frequency is approximately 261 Hz. Consult the text for even more details.

Musical notation shows which notes are to be played and their relative timing (half, quarter, or eighth). Figure 3 shows how the keys on the piano correspond to notes drawn in musical notation. The white keys are labeled as A, B, C, D, E, F, and G; but the black keys are denoted with “sharps” or “flats.” A sharp such as  $A^\sharp$  is one key number larger than A; a flat is one key lower, e.g.,  $A_4^\flat$  (A-flat) is key number 48.

<sup>2</sup>If you have little or no experience reading music, don’t be intimidated. Only a little music knowledge is needed to carry out this lab. On the other hand, the experience of working in an application area where you must quickly acquire new knowledge is a valuable one. Many real-world engineering problems have this flavor, especially in signal processing which has such a broad applicability in diverse areas such as geophysics, medicine, radar, speech, etc.

<sup>3</sup>In this lab, we are using the number 40 to represent middle C. This is somewhat arbitrary; for instance, the Musical Instrument Digital Interface (MIDI) standard represents middle C with the number 60.

The image shows a musical score for 'The Star-Spangled Banner' with two systems. The first system is for the first line of the song, and the second system is for the second line. The notation includes treble and bass staves with a key signature of one sharp (F#) and a 3/4 time signature. The first system has a forte (f) dynamic marking. The second system has a piano (p) dynamic marking. The notes are labeled with their names and durations: F-sharp, eighth note, half note, and quarter note. The second system also includes a list of notes with their corresponding frequencies and octave numbers: D4, E4, F#4, G4, A4, B4, C5, D5, C4 (middle-C), B3, A3, G3, F#3, E3, D3, C3, B2.

Thus, you can use the ratio  $2^{1/12}$  to calculate the frequency of notes anywhere on the piano keyboard. For example, the E-flat above middle-C (black key number 43) is 6 keys below A-440, so its frequency should be  $f_{43} = 440 \times 2^{-6/12} = 440 / \sqrt{2} \approx 311$  Hertz.

As done in Lab 3, it is often useful to represent the time-varying spectrum of a signal in terms of the *spectrogram* (see Chapter 3 in the text). A spectrogram is produced by estimating the frequency content in short time segments of the signal. The magnitude of the spectrum over individual segments is plotted as intensity or color on a two-dimensional plot versus frequency and time. Refresh your memory of Lab 3.

### 3.1 Debug coscos.m

1. Set a breakpoint to stop execution when an error occurs and jump into “Keyboard” mode;



2. Display the contents of important vectors while stopped;
3. Determine the size of all vectors by using either the **size()** function or the **whos** command; and,
4. Lastly, modify variables while in the “Keyboard” mode of the debugger.

```
function [xx,tt] = coscos( f1, f2, fs, dur )
% COSCOS multiply two sinusoids
%
t1 = 0:(1/fs):dur;
t2 = 0:(1/f2):dur;
cos1 = cos(2*pi*f1*t1);
cos2 = cos(2*pi*f2*t2);
xx = cos1 .* cos2;
tt = t1;
```

Show that you can use the debugger on the text file **coscos.m** to the TA for verification.

**Instructor Verification** (separate page)

### 3.2 Note Frequency Function

Write an M-file to produce a desired note, expressed in piano key number, for a given duration. Your M-file should be in the form of a function called **key2note.m**. Your function should have the following form:

```
function xx = key2note(X, keynum, dur, fsamp)
% KEY2NOTE Produce a sinusoidal waveform corresponding to a
% given piano key number
%
% usage: xx = key2note (X, keynum, dur)
%
% xx = the output sinusoidal waveform
% X = complex amplitude for the sinusoid, X = A*exp(j*phi).
% keynum = the piano keyboard number of the desired note
% dur = the duration (in seconds) of the output note
% fs = sampling frequency, use 11025 or 22050 Hz
%
tt = 0:(1/fsamp):dur;
freq = %<===== fill in this line
xx = real( X*exp(j*2*pi*freq*tt) );
```

For the “**freq** =” line, use the formulas given above to determine the frequency for a sinusoid in terms of its key number. You should start from a reference note (middle-C or A-440 is recommended) and solve for the frequency based on this reference. Notice that the **xx = real( )** line generates the actual sinusoid as the real part of a complex exponential at the proper frequency. Show the code and play a note, say key #47 (G4), for 2 seconds to the instructor for verification.

**Instructor Verification** (separate page)

### 3.3 Synthesize a Scale with Octaves

In a previous section you completed the **key2note.m** function which synthesizes the correct sinusoidal signal for a particular key number. Now, use that function to finish the following incomplete M-file that will play scales where each tone is the sum of two notes separated by an octave:

```

%--- play_scale_octave.m
%---
scale.keys = [ 40 42 44 45 47 49 51 52 40 44 47 44 40];
%--- NOTES: C D E F G A B C C E G E C
% key #40 is middle-C
%
scale.durations = 0.25 * ones(1,length(scale.keys));
fs = 11025; %-- or 22050 Hz
xx = zeros(1, ceil(sum(scale.durations)*fs+length(scale.keys)) );
n1 = 1;
for kk = 1:length(scale.keys)
    keynum = scale.keys(kk);

    tone =                                     %<===== FILL IN THIS LINE

    n2 = n1 + length(tone) - 1;
    xx(n1:n2) = xx(n1:n2) + tone;      %<=== Insert the note
    n1 = n2 + 1;
end
soundsc( xx, fs )

```

For the **tone** = line, generate the *two sinusoids*, one for **keynum**, the other for a key that is one octave higher. The sinusoids would be generated by making calls to the function **key2note()** written previously. It is important to point out that the code in **play\_scale\_octave.m** allocates a vector of zeros large enough to hold the entire scale then **inserts** each note into its proper place in the vector **xx**. Play the notes to the instructor.

**Instructor Verification** (separate page)

### 3.4 Spectrograms

Now, you will display the spectrogram of the scale synthesized in the previous section. Remember that the spectrogram displays an image that shows the *frequency* content of the synthesized *time* signal. Its horizontal axis is time and its vertical axis is frequency.

- (a) Generate the signal for the scale with **play scale octave.m**.
- (b) Use the function **spectrogram(xx,512,384,512,fs,'yaxis')**. Zoom in to see the progression of notes up the scale (**help zoom**), and also the notes one octave higher. In addition, identify the notes  $F$  (A-440) and  $A_5$  in your spectrogram.



CD-ROM  
Spectrogram  
& Sounds

**Instructor Verification** (separate page)

## 4 Optional Lab: Sinusoidal Chop Sticks

For this project, a very simple song is to be synthesized. Before starting the project, make sure that you have a working knowledge of the relationship between musical notes, piano key numbers and their frequencies (from the Pre-Lab).

### 4.1 Chop Sticks

“Chop Sticks” is one piano piece that everyone can play. A sequence of key-pairs is played, using the white keys in the octave starting at middle-C. The actual notes are given below:

The keys to be played are  $F_4$  and  $G_4$  (simultaneously) six times, then  $E_4$  and  $G_4$  six times, then  $D_4$  and  $B_4$  six times, then  $C_4$  (middle-C) and  $C_5$  four times, and finally,  $D_4$  and  $B_4$  once and  $E_4$  and  $A_4$  once. Then these 24 key-pairs make one cycle which can be repeated.

## 4.2 Synthesis of the Music Signal

The synthesis should be carried out with sinusoids. Two operations will be needed to create the song: addition of two sinusoids to handle the case where two notes are played simultaneously, and concatenation of signals to make the sequence of 24 notes for one cycle of Chop Sticks.

In the process of actually synthesizing the music, keep the following things in mind:

- The sampling frequency that will be used to play out the sound through the D-to-A system of the computer should be 11025 Hz.
- The time duration needed for each note-pair should be in the range of 150–250 msec. All note-pairs can have the same duration.
- There should be a short interval of silence between successive note-pairs; 50–100 msec should be sufficient to hear separate notes in time.
- To determine the frequency (in Hertz) for each note see Fig. 2 and the discussion of the well-tempered scale in the Pre-Lab.
- Set up your MATLAB synthesis program so that you can quickly generate one cycle of Chop Sticks and play it out for a listening.
- You can use `wavwrite` to save the synthesized result as a `wav` file.

## 4.3 Spectrogram of the Music

Musical notation describes how a song is composed of different frequencies and when they should be played. This representation is, in effect, a *time-frequency* representation of the signal that synthesizes the song. In MATLAB we can obtain a time-frequency representation from the synthesized signal itself by making a spectrogram with the MATLAB function `spectrogram()`, or `plotspec()`.

- Produce a spectrogram image of your synthesized music—one cycle of Chop Sticks. Since the spectrogram M-files will scale the frequency axis to run from zero to half the sampling frequency, it will be necessary to “zoom in” on the region where the notes are. The frequency region should be 0 to 600 Hz. Consult `help axis` or `help zoom`, or use the zoom tool in the MATLAB figure window, to display and print this region.
- In order to see all the notes in the spectrogram, you will probably have to adjust the window length from its default value of 256 in `spectrogram()`. Longer windows will give better frequency resolution, but if the window length is too long the spectrogram plot will become blurred along the time axis (why?).
- Try to verify on your spectrogram plot that you have the correct frequencies, and the correct number of notes.

## 4.4 Miscellaneous Comments

### 4.4.1 Music GUI

To aid your understanding of music and its connection to frequency content, a MATLAB GUI is available so that you can visualize the spectrogram along with musical notation. This GUI also has the capability to synthesize music from a list of notes, but these notes are given in “standard” musical notation, not key number. For more information, consult the help on `musicgui.m` which is part of the *SP-First* toolbox.



CD-ROM

MUSIC  
GUI

### 4.4.2 Musical Tweaks

The musical passage is likely to sound very artificial, because it is created from pure sinusoids. One annoying quality is a clicking sound that is caused by an abrupt beginning or end of a sinusoid. An envelope can be used to smooth the edges of the sinusoids. In addition, some harmonics could be added to enrich the sound. If you want to pursue further, these suggestions might lead you to even more impressive results.



#### 4.4.3 Envelope

One major improvement comes from using an “envelope,” where you multiply each pure tone signal by an envelope  $E(t)$  so that it fades in and out.

$$x(t) = E(t) \cos(2\pi f_{key} t + \varphi) \quad (2)$$

If an envelope is used it should “fade in” quickly and “fade out” more slowly. An envelope such as a half-cycle of a sine wave `sin( $\pi t / dur$ )` is simple to program, but it sounds poor.

A standard way to define the envelope function is to divide  $E(t)$  into four sections: attack (A), delay (D), sustain (S), and release (R). Together these are called ADSR. “Attack” is a quickly rising front edge, “delay” is a small short-duration drop, “sustain” is more or less constant and “release” drops quickly back to zero. Figure 4 shows a linear approximation to the ADSR profile. Consult help on `linspace()` or `interp1()` for functions that create linearly increasing and decreasing vectors.




Figure 4: ADSR profile for an envelope function  $E(t)$ .

**Lab #4**  
**ECE-2026 Fall-2013**  
**LAB SHEET**

*For each verification, be prepared to explain your answer and respond to other related questions that the lab TA's or professors might ask. Turn in the Lab Sheet at beginning of your next lab period.*

Name: \_\_\_\_\_

Date of Lab: \_\_\_\_\_

Part 3.1 Show that you can use the debugger on the text file **coscos.m**:

Verified: \_\_\_\_\_

Date/Time: \_\_\_\_\_

Part 3.2 Complete and demonstrate the function **key2note.m**:

Verified: \_\_\_\_\_

Date/Time: \_\_\_\_\_

Part 3.3 Complete and demonstrate the script file play scale **octave.m**:

Verified: \_\_\_\_\_

Date/Time: \_\_\_\_\_

Part 3.4 Demonstrate the spectrogram of the octave scale generated by **play\_scale\_octave.m**:

Verified: \_\_\_\_\_

Date/Time: \_\_\_\_\_

**Questions:**

Q4.1 Most pianos today have 88 keys, but one rare piano (Bösendorfer 290) has 97(!) keys. These extra keys are at the bass end of a regular piano keyboard. What is the lowest note frequency a Bösendorfer 290 can produce?

Q4.2 Suppose you want to break the world's record and make a piano with 106 keys. (You can easily do so with your computer!) You add the extra keys evenly at both ends of a regular piano keyboard, i.e., 9 extra keys above and below a regular keyboard, respectively. Use the function **key2note** that you have written in lab to synthesize the highest piano note so created. Throughout the lab exercises, you have been advised to use **fsamp = 11025** or **22050** (see 3.3). Produce the two output signals using these two numbers. Do they sound the same (don't forget to use the corresponding **fsamp** in **soundsc**) as intended? Try to explain or speculate what you have heard.